

Complete factorial designs

Session 5

MATH 80667A: Experimental Design and Statistical Methods
HEC Montréal

Outline

Factorial designs and interactions

Tests for two-way ANOVA

Factorial designs and interactions

Complete factorial designs?

Factorial design

study with multiple factors (subgroups)

Complete

Gather observations for every subgroup

Motivating example







Response:
retention of information
two hours after reading a story

Population:
children aged four

experimental factor 1:
ending (happy or sad)

experimental factor 2:
complexity (easy, average or hard).

Setup of design

complexity	happy	sad
complicated		
average		
easy		

Efficiency of factorial design

**Cast problem
as a series of one-way ANOVA
vs simultaneous estimation**

**Factorial designs requires
fewer overall observations**

Can study interactions

Interaction

Definition: when the effect of one factor depends on the levels of another factor.


Effect together
 \neq
sum of individual effects

Interaction or profile plot


**Graphical display:
plot sample mean per category**

**with uncertainty measure
(1 std. error for mean
confidence interval, etc.)**

Interaction plots and parallel lines




complexity — average — complicated — easy



complexity — average — complicated — easy

Interaction plots for 2 by 2 designs



Cell means for 2 by 2 designs

	b1	b2
<i>a1</i>	5	5
<i>a2</i>	5	5

	b1	b2
<i>a1</i>	10	10
<i>a2</i>	5	5

	b1	b2
<i>a1</i>	5	10
<i>a2</i>	5	10

	b1	b2
<i>a1</i>	5	10
<i>a2</i>	10	15

	b1	b2
<i>a1</i>	5	10
<i>a2</i>	10	5

	b1	b2
<i>a1</i>	10	13
<i>a2</i>	5	2

	b1	b2
<i>a1</i>	2	12
<i>a2</i>	5	9

	b1	b2
<i>a1</i>	10	18
<i>a2</i>	5	7

Example 1 : loans versus credit

Sharma, Tully, and Cryder (2021)

Supplementary study 5 consists of a 2×2 between-subject ANOVA with factors

- debt type (debttype), either "loan" or "credit"
- purchase type, either discretionary or not (need)

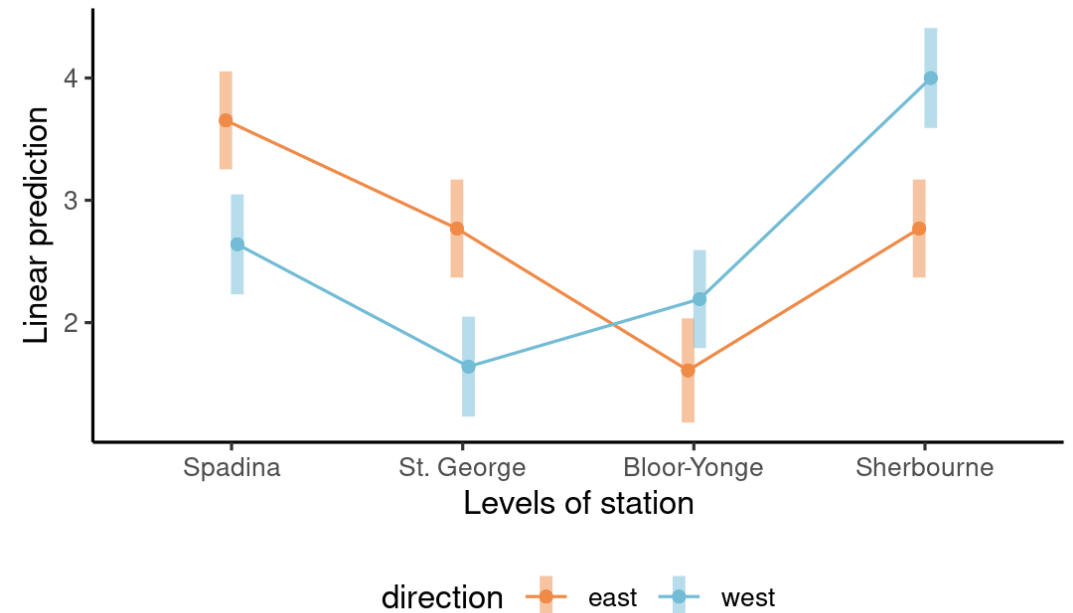


No evidence of interaction

Example 2 - psychological distance

Maglio and Polman (2014) Study 1
uses a 4×2 between-subject
ANOVA with factors

- subway station, one of Spadina, St. George, Bloor-Yonge and Sherbourne
- direction of travel, either east or west



Clear evidence of interaction
(symmetry?)

Tests for two-way ANOVA

Analysis of variance = regression

An analysis of variance model is simply a **linear regression** with categorical covariate(s).

- Typically, the parametrization is chosen so that parameters reflect differences to the global mean (sum-to-zero parametrization).
- The full model includes interactions between all combinations of factors
 - one average for each subcategory
 - one-way ANOVA!

Formulation of the two-way ANOVA

Two factors: A (complexity) and B (ending) with $n_a = 3$ and $n_b = 2$ levels, and their interaction.

Write the average response Y_{ijr} of the r th measurement in group (a_i, b_j) as

$$\begin{array}{ccc} \mathbf{E}(Y_{ijr}) & = & \mu_{ij} \\ \text{average response} & & \text{subgroup mean} \end{array}$$

where Y_{ijr} are independent observations with a common std. deviation σ .

- We estimate μ_{ij} by the sample mean of the subgroup (i, j) , say $\hat{\mu}_{ij}$.
- The fitted values are $\hat{y}_{ijr} = \hat{\mu}_{ij}$.

One average for each subgroup

B ending A complexity	b_1 (happy)	b_2 (sad)	row mean
a_1 (complicated)	μ_{11}	μ_{12}	$\mu_{1.}$
a_2 (average)	μ_{21}	μ_{22}	$\mu_{2.}$
a_3 (easy)	μ_{31}	μ_{32}	$\mu_{3.}$
<i>column mean</i>	$\mu_{.1}$	$\mu_{.2}$	μ

Row, column and overall average

- Mean of A_i (average of row i):

$$\mu_{i.} = \frac{\mu_{i1} + \cdots + \mu_{in_b}}{n_b}$$

- Mean of B_j (average of column j):

$$\mu_{.j} = \frac{\mu_{1j} + \cdots + \mu_{n_a j}}{n_a}$$

- Overall average:

$$\mu = \frac{\sum_{i=1}^{n_a} \sum_{j=1}^{n_b} \mu_{ij}}{n_a n_b}$$

- Row, column and overall averages are **equiweighted** combinations of the cell means μ_{ij} .
- Estimates are obtained by replacing μ_{ij} in formulas by subgroup sample mean.

Vocabulary of effects

- **simple effects**: difference between levels of one in a fixed combination of others (change in difficulty for happy ending)
- **main effects**: differences relative to average for each condition of a factor (happy vs sad ending)
- **interaction effects**: when simple effects differ depending on levels of another factor

Main effects

Main effects are comparisons between row or column averages

Obtained by *marginalization*, i.e., averaging over the other dimension.

Main effects are not of interest if there is an interaction.

	happy sad	
column means	$\mu_{.1}$	$\mu_{.2}$

complexity	row means
complicated	$\mu_1.$
average	$\mu_2.$
easy	$\mu_3.$

Simple effects

Simple effects are comparisons between cell averages within a given row or column

	happy sad	
means (easy)	μ_{13}	μ_{23}

complexity	mean (happy)
complicated	μ_{11}
average	μ_{21}
easy	μ_{31}

Contrasts

We collapse categories to obtain a one-way ANOVA with categories A (complexity) and B (ending).

Q: How would you write the weights for contrasts for testing the

- main effect of A : complicated vs average, or complicated vs easy.
- main effect of B : happy vs sad.
- interaction A and B : difference between complicated and average, for happy versus sad?

The order of the categories is $(a_1, b_1), (a_1, b_2), \dots, (a_3, b_2)$.

Contrasts

Suppose the order of the coefficients is factor A (complexity, 3 levels, complicated/average/easy) and factor B (ending, 2 levels, happy/sad).

test	μ_{11}	μ_{12}	μ_{21}	μ_{22}	μ_{31}	μ_{32}
main effect A (complicated vs average)	1	1	-1	-1	0	0
main effect A (complicated vs easy)	1	1	0	0	-1	-1
main effect B (happy vs sad)	1	-1	1	-1	1	-1
interaction AB (comp. vs av, happy vs sad)	1	-1	-1	1	0	0
interaction AB (comp. vs easy, happy vs sad)	1	-1	0	0	-1	1

Global hypothesis tests

Main effect of factor A

$\mathcal{H}_0: \mu_{1.} = \dots = \mu_{n_a.}$ vs \mathcal{H}_a : at least two marginal means of A are different

Main effect of factor B

$\mathcal{H}_0: \mu_{.1} = \dots = \mu_{.n_b}$ vs \mathcal{H}_a : at least two marginal means of B are different

Interaction

$\mathcal{H}_0: \mu_{ij} = \mu_{i.} + \mu_{.j}$ (sum of main effects) vs \mathcal{H}_a : effect is not a combination of row/column effect.

Comparing nested models

Rather than present the specifics of ANOVA, we consider a general hypothesis testing framework which is more widely applicable.

We compare two competing models

- the alternative or full model \mathcal{H}_a
- the simpler null model \mathcal{H}_0 , which imposes ν restrictions on the full model

Intuition behind F -test for ANOVA

The more complex fits better (it is necessarily more flexible), but requires estimation of more parameters.

- Test compares the goodness-of-fit and attempts to determine what is the improvement that would occur by chance, if the null model was correct, given ν additional parameters.

Testing linear restrictions in linear models

If the alternative model has p parameters for the mean, and we impose ν linear restrictions under the null hypothesis to the model estimated based on n independent observations, the test statistic is

$$F = \frac{(\text{RSS}_0 - \text{RSS}_a)/\nu}{\text{RSS}_a/(n - p)}$$

- The numerator is the difference in residuals sum of squares, denoted RSS , from models fitted under \mathcal{H}_0 and \mathcal{H}_a , divided by degrees of freedom ν .
- The denominator is an estimator of the variance, obtained under \mathcal{H}_a (termed mean squared error of residuals)
- The benchmark for tests in linear models is Fisher's $F(\nu, n - p)$.

Analysis of deviance

For other generalized linear models with parameters $\boldsymbol{\theta}$, we proceed similarly, but we use the log *likelihood* function, $\ell(\boldsymbol{\theta})$ as goodness-of-fit measure.

- The higher the (log) likelihood, the better the fit.
- Obtain parameter estimates $\hat{\boldsymbol{\theta}}_0$ under the null hypothesis and $\hat{\boldsymbol{\theta}}_a$ under the alternative by maximum likelihood estimation.
- Consider the likelihood ratio statistic

$$R = 2\{\ell(\hat{\boldsymbol{\theta}}_a) - \ell(\hat{\boldsymbol{\theta}}_0)\}$$

- Under regularity conditions, we compare R to a chi-square distribution with ν degrees of freedom, χ^2_ν .

Analysis of variance table

term	degrees of freedom	mean square	F
A	$n_a - 1$	$MS_A = SS_A / (n_a - 1)$	MS_A / MS_{res}
B	$n_b - 1$	$MS_B = SS_B / (n_b - 1)$	MS_B / MS_{res}
AB	$(n_a - 1)(n_b - 1)$	$MS_{AB} = SS_{AB} / \{(n_a - 1)(n_b - 1)\}$	$MS_{AB} / MS_{\text{res}}$
residuals	$n - n_a n_b$	$MS_{\text{resid}} = RSS_a / (n - ab)$	
total	$n - 1$		

Read the table backward (starting with the interaction).

- If there is a significant interaction, the main effects are **not** of interest and potentially misleading.

Intuition behind degrees of freedom

The model always includes an overall average μ . There are

- $n_a - 1$ free row means since $n_a\mu = \mu_{1.} + \dots + \mu_{n_a.}$
- $n_b - 1$ free column means as $n_b\mu = \mu_{.1} + \dots + \mu_{.n_b}$
- $n_an_b - (n_a - 1) - (n_b - 1) - 1$ interaction terms

<i>B</i> ending <i>A</i> complexity	<i>b</i> ₁ (happy)	<i>b</i> ₂ (sad)	row mean
<i>a</i> ₁ (complicated)	μ_{11}	X	$\mu_{1.}$
<i>a</i> ₂ (average)	μ_{21}	X	$\mu_{2.}$
<i>a</i> ₃ (easy)	X	X	X
column mean	$\mu_{.1}$	X	μ

Terms with X are fully determined by row/column/total averages

Example 1

The interaction plot suggested that the two-way interaction wasn't significant. The F test confirms this.

There is a significant main effect of both `purchase` and `debttype`.

term	SS	df	F stat	p-value
purchase	752.3	1	98.21	0.000
debttype	92.2	1	12.04	0.001
purchase:debttype	13.7	1	1.79	0.182
Residuals	11467.4	1497		

Example 2

There is a significant interaction between `station` and `direction`, so follow-up by looking at simple effects or contrasts.

The tests for the main effects are not of interest! Disregard other entries of the ANOVA table

term	SS	df	F stat	p-value
station	75.2	3	23.35	0.000
direction	0.4	1	0.38	0.541
station:direction	52.4	3	16.28	0.000
Residuals	208.2	194		


Main effects for Example 1

We consider differences between debt type labels.

Participants are more likely to consider the offer if it is branded as `credit` than `loan`. The difference is roughly 0.5 (on a Likert scale from 1 to 9).

```
## $emmeans
##   debttype emmean      SE    df lower.CL upper.CL
##   credit      5.12 0.101 1497      4.93      5.32
##   loan        4.63 0.101 1497      4.43      4.83
##
## Results are averaged over the levels of: purchase
## Confidence level used: 0.95
##
## $contrasts
##   contrast      estimate      SE    df t.ratio p.value
##   credit - loan      0.496 0.143 1497      3.469  0.0005
##
## Results are averaged over the levels of: purchase
```

Toronto subway station



Simplified depiction of the Toronto metro stations used in the experiment, based on work by Craftworker on Wikipedia, distributed under CC-BY-SA 4.0.


Reparametrization for Example 2

Set `stdist` as $-2, -1, +1, +2$ to indicate station distance, with negative signs indicating stations in opposite direction of travel

The ANOVA table for the reparametrized models shows no evidence against the null of symmetry (interaction).

term	SS	df	F stat	p-value
stdist	121.9	3	37.86	0.000
direction	0.4	1	0.35	0.554
stdist:direction	5.7	3	1.77	0.154
Residuals	208.2	194		

Interaction plot for reformed data



Custom contrasts for Example 2

We are interested in testing the perception of distance, by looking at

$$\mathcal{H}_0 : \mu_{-1} = \mu_{+1}, \mu_{-2} = \mu_{+2}.$$

```
mod3 <- lm(distance ~ stdist * direction, data = MP14_S1)
(emm <- emmeans(mod3, specs = "stdist"))
  # order is -2, -1, 1, 2
contrasts <- emm |> contrast(
  list("two dist" = c(-1, 0, 0, 1),
       "one dist" = c(0, -1, 1, 0)))
contrasts # print pairwise contrasts
test(contrasts, joint = TRUE)
```

Estimated marginal means and contrasts

Strong evidence of differences in perceived distance depending on direction of travel.

```
##      stdist emmean      SE   df lower.CL upper.CL
##    -2          3.83 0.145 194      3.54      4.11
##    -1          2.48 0.144 194      2.20      2.76
##    +1          1.62 0.150 194      1.33      1.92
##    +2          2.70 0.145 194      2.42      2.99
##
## Results are averaged over the levels of: direction
## Confidence level used: 0.95
```

```
##      contrast estimate      SE   df t.ratio p.value
##    two dist      -1.122 0.205 194   -5.470  <.0001
##    one dist      -0.856 0.207 194   -4.129  0.0001
##
## Results are averaged over the levels of: direction
```

```
##      df1 df2 F.ratio p.value
##        2 194   23.485  <.0001
```